

Tidy data & dplyr

Lecture 06

Dr. Colin Rundel



Tidy data

country	year	cases	population
Afghanistan	1999	7645	19987071
Afghanistan	2000	8666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	214258	1272915272
China	2000	216766	128042583

variables

country	year	cases	population
Afghanistan	1999	7645	19987071
Afghanistan	2000	8666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	214258	1272915272
China	2000	216766	128042583

observations

country	year	cases	population
Afghanistan	1999	7645	19987071
Afghanistan	2000	8666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	214258	1272915272
China	2000	216766	128042583

values

Tidy vs Untidy

Happy families are all alike; every unhappy family is unhappy in its own way
— Leo Tolstoy, Anna Karenina

```
# A tibble: 317 × 7
  artist      track      date.ent...1  wk1  wk2  wk3  wk4
  <chr>      <chr>      <date>      <dbl> <dbl> <dbl> <dbl>
1 2 Pac      Baby Don't Cry (Keep... 2000-02-26    87    82    72    77
2 2Ge+her    The Hardest Part Of ... 2000-09-02    91    87    92    NA
3 3 Doors Down Kryptonite      2000-04-08    81    70    68    67
4 3 Doors Down Loser            2000-10-21    76    76    72    69
5 504 Boyz    Wobble Wobble    2000-04-15    57    34    25    17
6 98^0        Give Me Just One Nig... 2000-08-19    51    39    34    26
7 A*Teens     Dancing Queen    2000-07-08    97    97    96    95
8 Aaliyah     I Don't Wanna    2000-01-29    84    62    51    41
9 Aaliyah     Try Again        2000-03-18    59    53    38    28
10 Adams     Yolanda Open My Heart 2000-08-26    76    76    74    60
```

Is the above data set tidy?

More tidy vs untidy

Is the following data tidy?

List of 3

```
$ :List of 8
..$ name      : chr "Luke Skywalker"
..$ height    : chr "172"
..$ mass      : chr "77"
..$ hair_color: chr "blond"
..$ skin_color: chr "fair"
..$ eye_color : chr "blue"
..$ birth_year: chr "19BBY"
..$ gender    : chr "male"
$ :List of 8
..$ name      : chr "C-3PO"
..$ height    : chr "167"
```

List of 3

```
$ :List of 8
..$ name      : chr "Darth Vader"
..$ height    : chr "202"
..$ mass      : chr "136"
..$ hair_color: chr "none"
..$ skin_color: chr "white"
..$ eye_color : chr "yellow"
..$ birth_year: chr "41.9BBY"
..$ gender    : chr "male"
$ :List of 8
..$ name      : chr "Leia Organa"
..$ height    : chr "150"
```



Modern data frames

The tidyverse includes the tibble package that extends data frames to be a bit more modern. The core features of tibbles is to have a nicer printing method as well as being “surly” and “lazy”.

```
1 library(tibble)
```

```
1 iris
```

	Sepal.Length	Sepal.Width	Petal.Length
1	5.1	3.5	1.4
2	4.9	3.0	1.4
3	4.7	3.2	1.3
4	4.6	3.1	1.5
5	5.0	3.6	1.4
6	5.4	3.9	1.7
7	4.6	3.4	1.4
8	5.0	3.4	1.5
9	4.4	2.9	1.4
10	4.9	3.1	1.5
11	5.4	3.7	1.5
12	4.8	3.4	1.6
13	4.8	3.0	1.4
14	4.3	3.0	1.1
15	5.8	4.0	1.2
16	5.7	4.4	1.5

```
1 (tbl_iris = as_tibble(iris))
```

```
# A tibble: 150 × 5
  Sepal.Length Sepal.Wi...1 Petal...2 Petal...3 Species
      <dbl>         <dbl>      <dbl>      <dbl> <fct>
1         5.1         3.5        1.4        0.2 setosa
2         4.9         3.0        1.4        0.2 setosa
3         4.7         3.2        1.3        0.2 setosa
4         4.6         3.1        1.5        0.2 setosa
5          5          3.6        1.4        0.2 setosa
6         5.4         3.9        1.7        0.4 setosa
7         4.6         3.4        1.4        0.3 setosa
8          5          3.4        1.5        0.2 setosa
9         4.4         2.9        1.4        0.2 setosa
10        4.9         3.1        1.5        0.1 setosa
# ... with 140 more rows, and abbreviated variable
#   names 1Sepal.Width, 2Petal.Length,
#   3Petal.Width
```

Tibbles are lazy

By default, subsetting tibbles always results in another tibble (`$` or `[[` can still be used to subset for a specific column). I.e. tibble subsets are always preserving and therefore type consistent.

```
1 tbl_iris[1,]
```

```
# A tibble: 1 × 5
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
	<dbl>	<dbl>	<dbl>	<dbl>	<fct>
1	5.1	3.5	1.4	0.2	setosa

More laziness - partial matching

Tibbles do not use partial matching when the `$` operator is used.

```
1 head( iris$Species )
```

```
[1] setosa setosa setosa setosa setosa  
setosa
```

```
Levels: setosa versicolor virginica
```

```
1 head( iris$Sp )
```

```
[1] setosa setosa setosa setosa setosa  
setosa
```

```
Levels: setosa versicolor virginica
```

```
1 head( tbl_iris$Species )
```

```
[1] setosa setosa setosa setosa setosa  
setosa
```

```
Levels: setosa versicolor virginica
```

```
1 head( tbl_iris$Sp )
```

```
NULL
```

More laziness - stringsAsFactors

Tibbles also have always had `stringsAsFactors = FALSE` as default behavior.

```
1 (t = tibble(  
2   x = 1:3,  
3   y = c("A", "B", "C"),  
4   z = factor(c("X", "Y", "Z"))  
5 ))
```

```
# A tibble: 3 × 3
```

	x	y	z
	<int>	<chr>	<fct>
1	1	A	X
2	2	B	Y
3	3	C	Z

Tibbles and length coercion

Only vectors with length 1 will undergo length coercion - everything else will throw an error.

```
1 data.frame(x = 1:4, y = 1)
```

	x	y
1	1	1
2	2	1
3	3	1
4	4	1

```
1 tibble(x = 1:4, y = 1)
```

```
# A tibble: 4 × 2
      x     y
  <int> <dbl>
1     1     1
2     2     1
3     3     1
4     4     1
```

```
1 data.frame(x = 1:4, y = 1:2)
```

	x	y
1	1	1
2	2	2
3	3	1
4	4	2

```
1 tibble(x = 1:4, y = 1:2)
```

```
Error:
! Tibble columns must have compatible
sizes.
• Size 4: Existing data.
• Size 2: Column `y`.
i Only values of size one are recycled.
```

Tibbles and S3

```
1 t = tibble(  
2   x = 1:3,  
3   y = c("A", "B", "C")  
4 )  
5  
6 class(t)
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

```
1 d = data.frame(  
2   x = 1:3,  
3   y = c("A", "B", "C")  
4 )  
5  
6 class(d)
```

```
[1] "data.frame"
```

```
1 methods(class="tbl_df")
```

```
[1] [          [[          [[<-          [<-          $  
[6] $<-          as.data.frame coerce          initialize names<-  
[11] Ops          row.names<-  show          slotsFromS3  str  
[16] tbl_sum
```

see '?methods' for accessing help and source code

```
1 methods(class="tbl")
```

```
[1] [[<-          [<-          $<-          coerce          format  
[6] glimpse          initialize Ops          print          show  
[11] slotsFromS3 tbl_sum
```

see '?methods' for accessing help and source code

Supporting tibbles?

```
1 d = tibble(  
2   x = rnorm(100),  
3   y = 3 + x + rnorm(100, sd = 0.1)  
4 )
```

```
1 lm(y~x, data = d)
```

Call:

```
lm(formula = y ~ x, data = d)
```

Coefficients:

(Intercept)	x
2.9947	0.9954

Why did this work?

magrittr



What is a pipe

In software engineering, a pipeline consists of a chain of processing elements (processes, threads, coroutines, functions, etc.), arranged so that the output of each element is the input of the next; - [Wikipedia - Pipeline \(software\)](#)

Magrittr's pipe is a new infix operator that allows us to link two functions together in a way that is readable from left to right.

The two code examples below are equivalent,

```
1 f(g(x=1, y=2), n=2)
```

```
1 g(x=1, y=2) %>% f(n=2)
```

Readability

Consider the following sequence of actions that describe the process of getting to campus in the morning:

I need to find my key, then unlock my car, then start my car, then drive to school, then park.

Expressed as a set of nested functions in R pseudocode this would look like:

```
1 park(drive(start_car(find("keys")), to="campus"))
```

Writing it out using pipes give it a more natural (and easier to read) structure:

```
1 find("keys") %>%  
2   start_car() %>%  
3   drive(to="campus") %>%  
4   park()
```

Approaches

All of the following are fine, it comes down to personal preference:

Nested:

```
1 h( g( f(x), y=1 ), z=1 )
```

Piped:

```
1 f(x) %>%  
2   g(y=1) %>%  
3   h(z=1)
```

Intermediate:

```
1 res = f(x)  
2 res = g(res, y=1)  
3 res = h(res, z=1)
```

What about other arguments?

Sometimes we want to send our results to an function argument other than first one or we want to use the previous result for multiple arguments. In these cases we can refer to the previous result using `..`

```
1 data.frame(a = 1:3, b = 3:1) %>% lm(a~b, data=.)
```

Call:

```
lm(formula = a ~ b, data = .)
```

Coefficients:

(Intercept)	b
4	-1

```
1 data.frame(a = 1:3, b = 3:1) %>% .[[1]]
```

```
[1] 1 2 3
```

```
1 data.frame(a = 1:3, b = 3:1) %>% .[[length(.)]]
```

```
[1] 3 2 1
```

The base R pipe

As of R v4.1.0 a pipe operator has been added to the base language in R, it is implemented as `|>`.

```
1 1:10 |> cumsum()
```

```
[1] 1 3 6 10 15 21 28 36 45 55
```


```
1 1:10 |> cumsum() |> mean()
```

```
[1] 22
```

The current version of RStudio on the departmental servers is v4.1 so you are welcome to try it out.

Base R pipe considerations:

- Depending on an R version ≥ 4.1 is a harder dependency than depending on the magrittr package
- `|>` does not support using `.` to pass returned values to other argument positions
- `|>` will likely have less overhead than `%>%` but the difference is unlikely to matter in practice
- `|>` supports an equivalent to `.` using `_` as of R v4.2



A Grammar of Data Manipulation

dplyr is based on the concepts of functions as verbs that manipulate data frames.

Core single data frame functions / verbs:

- `filter()` / `slice()`: pick rows based on criteria
- `select()` / `rename()`: select columns by name
- `pull()`: grab a column as a vector
- `arrange()`: reorder rows
- `mutate()` / `transmute()`: create or modify columns
- `distinct()`: filter for unique rows
- `summarise()` / `count()`: reduce variables to values
- `group_by()` / `ungroup()`: modify other verbs to act on subsets
- `relocate()`: change column order
- ... (many more)

dplyr heuristics

1. First argument is *always* a data frame
2. Subsequent arguments say what to do with that data frame
3. *Always* return a data frame
4. Don't modify in place
5. Magic via lazy evaluation

Example Data

We will demonstrate dplyr's functionality using the nycflights13 data.

```
1 library(dplyr)
2 library(nycflights13)
```

```
1 flights
```

```
# A tibble: 336,776 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	1	1	517	515	2	830	819	11
2	2013	1	1	533	529	4	850	830	20
3	2013	1	1	542	540	2	923	850	33
4	2013	1	1	544	545	-1	1004	1022	-18
5	2013	1	1	554	600	-6	812	837	-25
6	2013	1	1	554	558	-4	740	728	12
7	2013	1	1	555	600	-5	913	854	19
8	2013	1	1	557	600	-3	709	723	-14
9	2013	1	1	557	600	-3	838	846	-8
10	2013	1	1	558	600	-2	753	745	8

```
# ... with 336,766 more rows, 10 more variables: carrier <chr>,
```

filter() - March flights

```
1 flights %>% filter(month == 3)
```

```
# A tibble: 28,834 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	3	1	4	2159	125	318	56	142
2	2013	3	1	50	2358	52	526	438	48
3	2013	3	1	117	2245	152	223	2354	149
4	2013	3	1	454	500	-6	633	648	-15
5	2013	3	1	505	515	-10	746	810	-24
6	2013	3	1	521	530	-9	813	827	-14
7	2013	3	1	537	540	-3	856	850	6
8	2013	3	1	541	545	-4	1014	1023	-9
9	2013	3	1	549	600	-11	639	703	-24
10	2013	3	1	550	600	10	717	801	11

filter() - Flights in the first 7 days of March

```
1 flights %>% filter(month == 3, day <= 7)
```

```
# A tibble: 6,530 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	3	1	4	2159	125	318	56	142
2	2013	3	1	50	2358	52	526	438	48
3	2013	3	1	117	2245	152	223	2354	149
4	2013	3	1	454	500	-6	633	648	-15
5	2013	3	1	505	515	-10	746	810	-24
6	2013	3	1	521	530	-9	813	827	-14
7	2013	3	1	537	540	-3	856	850	6
8	2013	3	1	541	545	-4	1014	1023	-9
9	2013	3	1	549	600	-11	639	703	-24
10	2013	3	1	550	600	10	717	801	11

filter() - Flights to LAX or JFK in March

```
1 flights %>% filter(dest == "LAX" | dest == "JFK", month==3)
```

```
# A tibble: 1,178 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	3	1	607	610	-3	832	925	-53
2	2013	3	1	629	632	-3	844	952	-68
3	2013	3	1	657	700	-3	953	1034	-41
4	2013	3	1	714	715	-1	939	1037	-58
5	2013	3	1	716	710	6	958	1035	-37
6	2013	3	1	727	730	-3	1007	1100	-53
7	2013	3	1	836	840	-4	1111	1157	-46
8	2013	3	1	857	900	-3	1202	1221	-19
9	2013	3	1	903	900	3	1157	1220	-23
10	2013	3	1	000	001	00	1150	1151	1

slice() - First 10 flights

```
1 flights %>% slice(1:10)
```

```
# A tibble: 10 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	1	1	517	515	2	830	819	11
2	2013	1	1	533	529	4	850	830	20
3	2013	1	1	542	540	2	923	850	33
4	2013	1	1	544	545	-1	1004	1022	-18
5	2013	1	1	554	600	-6	812	837	-25
6	2013	1	1	554	558	-4	740	728	12
7	2013	1	1	555	600	-5	913	854	19
8	2013	1	1	557	600	-3	709	723	-14
9	2013	1	1	557	600	-3	838	846	-8
10	2013	1	1	558	600	-2	752	745	7

slice() - Last 5 flights

```
1 flights %>% slice((n()-4):n())
```

```
# A tibble: 5 × 19
```

	year	month	day	dep_t... ¹	sched... ²	dep_d... ³	arr_t... ⁴	sched... ⁵	arr_d... ⁶	carrier
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<chr>
1	2013	9	30	NA	1455	NA	NA	1634	NA	9E
2	2013	9	30	NA	2200	NA	NA	2312	NA	9E
3	2013	9	30	NA	1210	NA	NA	1330	NA	MQ
4	2013	9	30	NA	1159	NA	NA	1344	NA	MQ
5	2013	9	30	NA	840	NA	NA	1020	NA	MQ

```
# ... with 9 more variables: flight <int>, tailnum <chr>, origin <chr>,  
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>, and abbreviated variable names 1dep_time,  
#   2sched_dep_time, 3dep_delay, 4arr_time, 5sched_arr_time, 6arr_delay
```

slice_tail() - Last 5 flights

```
1 flights %>% slice_tail(n = 5)
```

```
# A tibble: 5 × 19
```

	year	month	day	dep_t... ¹	sched... ²	dep_d... ³	arr_t... ⁴	sched... ⁵	arr_d... ⁶	carrier
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<chr>
1	2013	9	30	NA	1455	NA	NA	1634	NA	9E
2	2013	9	30	NA	2200	NA	NA	2312	NA	9E
3	2013	9	30	NA	1210	NA	NA	1330	NA	MQ
4	2013	9	30	NA	1159	NA	NA	1344	NA	MQ
5	2013	9	30	NA	840	NA	NA	1020	NA	MQ

```
# ... with 9 more variables: flight <int>, tailnum <chr>, origin <chr>,  
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>, and abbreviated variable names 1dep_time,  
#   2sched_dep_time, 3dep_delay, 4arr_time, 5sched_arr_time, 6arr_delay
```


select() - Individual Columns

```
1 flights %>% select(year, month, day)
```

```
# A tibble: 336,776 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

select() - Exclude Columns

```
1 flights %>% select(-year, -month, -day)
```

A tibble: 336,776 × 16

	dep_time	sched_... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵	carrier	flight	tailnum
	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<chr>	<int>	<chr>
1	517	515	2	830	819	11	UA	1545	N14228
2	533	529	4	850	830	20	UA	1714	N24211
3	542	540	2	923	850	33	AA	1141	N619AA
4	544	545	-1	1004	1022	-18	B6	725	N804JB
5	554	600	-6	812	837	-25	DL	461	N668DN
6	554	558	-4	740	728	12	UA	1696	N39463
7	555	600	-5	913	854	19	B6	507	N516JB
8	557	600	-3	709	723	-14	EV	5708	N829AS
9	557	600	-3	838	846	-8	B6	79	N593JB
10	558	600	2	752	715	0	AA	201	N287AA

select() - Ranges

```
1 flights %>% select(year:day)
```

```
# A tibble: 336,776 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

select() - Exclusion Ranges

```
1 flights %>% select(-(year:day))
```

A tibble: 336,776 × 16

	dep_time	sched_... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵	carrier	flight	tailnum
	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<chr>	<int>	<chr>
1	517	515	2	830	819	11	UA	1545	N14228
2	533	529	4	850	830	20	UA	1714	N24211
3	542	540	2	923	850	33	AA	1141	N619AA
4	544	545	-1	1004	1022	-18	B6	725	N804JB
5	554	600	-6	812	837	-25	DL	461	N668DN
6	554	558	-4	740	728	12	UA	1696	N39463
7	555	600	-5	913	854	19	B6	507	N516JB
8	557	600	-3	709	723	-14	EV	5708	N829AS
9	557	600	-3	838	846	-8	B6	79	N593JB
10	558	600	2	752	715	0	AA	201	N287AA

select() - Matching contains()

```
1 flights %>% select(contains("dep"),
2                   contains("arr"))
```

A tibble: 336,776 × 7

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_t... ¹	arr_d... ²	carrier
	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<chr>
1	517	515	2	830	819	11	UA
2	533	529	4	850	830	20	UA
3	542	540	2	923	850	33	AA
4	544	545	-1	1004	1022	-18	B6
5	554	600	-6	812	837	-25	DL
6	554	558	-4	740	728	12	UA
7	555	600	-5	913	854	19	B6
8	557	600	-3	709	723	-14	EV
9	557	600	-3	838	846	-8	B6
10	550	600	2	752	715	0	AA

select() - Matching starts_with()

```
1 flights %>% select(starts_with("dep"),
2                     starts_with("arr"))
```

A tibble: 336,776 × 4

	dep_time	dep_delay	arr_time	arr_delay
	<int>	<dbl>	<int>	<dbl>
1	517	2	830	11
2	533	4	850	20
3	542	2	923	33
4	544	-1	1004	-18
5	554	-6	812	-25
6	554	-4	740	12
7	555	-5	913	19
8	557	-3	709	-14
9	557	-3	838	-8
10	558	2	752	0

Other helpers provide by tidyselect:

select() + where() - Get numeric columns

```
1 flights %>% select(where(is.numeric))
```

```
# A tibble: 336,776 × 14
```

	year	month	day	dep_t... ¹	sched... ²	dep_d... ³	arr_t... ⁴	sched... ⁵	arr_d... ⁶	flight
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830	819	11	1545
2	2013	1	1	533	529	4	850	830	20	1714
3	2013	1	1	542	540	2	923	850	33	1141
4	2013	1	1	544	545	-1	1004	1022	-18	725
5	2013	1	1	554	600	-6	812	837	-25	461
6	2013	1	1	554	558	-4	740	728	12	1696
7	2013	1	1	555	600	-5	913	854	19	507
8	2013	1	1	557	600	-3	709	723	-14	5708
9	2013	1	1	557	600	-3	838	846	-8	79
10	2013	1	1	558	600	-2	753	745	8	301

```
# ... with 336,766 more rows, 4 more variables: air_time <dbl>,
```

select() + where() - Get non-numeric columns

```
1 flights %>% select(where(function(x) !is.numeric(x)))
```

```
# A tibble: 336,776 × 5
```

	carrier	tailnum	origin	dest	time_hour
	<chr>	<chr>	<chr>	<chr>	<dtm>
1	UA	N14228	EWR	IAH	2013-01-01 05:00:00
2	UA	N24211	LGA	IAH	2013-01-01 05:00:00
3	AA	N619AA	JFK	MIA	2013-01-01 05:00:00
4	B6	N804JB	JFK	BQN	2013-01-01 05:00:00
5	DL	N668DN	LGA	ATL	2013-01-01 06:00:00
6	UA	N39463	EWR	ORD	2013-01-01 05:00:00
7	B6	N516JB	EWR	FLL	2013-01-01 06:00:00
8	EV	N829AS	LGA	IAD	2013-01-01 06:00:00
9	B6	N593JB	JFK	MCO	2013-01-01 06:00:00
10	AA	N3ALAA	LGA	ORD	2013-01-01 06:00:00

```
# ... with 336,766 more rows
```


relocate - to the front

```
1 flights %>% relocate(carrier, origin, dest)
```

```
# A tibble: 336,776 × 19
```

	carrier	origin	dest	year	month	day	dep_time	sched_... ¹	dep_d... ²	arr_t... ³
	<chr>	<chr>	<chr>	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	UA	EWR	IAH	2013	1	1	517	515	2	830
2	UA	LGA	IAH	2013	1	1	533	529	4	850
3	AA	JFK	MIA	2013	1	1	542	540	2	923
4	B6	JFK	BQN	2013	1	1	544	545	-1	1004
5	DL	LGA	ATL	2013	1	1	554	600	-6	812
6	UA	EWR	ORD	2013	1	1	554	558	-4	740
7	B6	EWR	FLL	2013	1	1	555	600	-5	913
8	EV	LGA	IAD	2013	1	1	557	600	-3	709
9	B6	JFK	MCO	2013	1	1	557	600	-3	838
10	AA	LGA	ORD	2013	1	1	558	600	2	752

relocate - to the end

```
1 flights %>%
2   relocate(year, month, day, .after = last_col())
```

A tibble: 336,776 × 19

	dep_time	sched_... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵	carrier	flight	tailnum
	<int>	<int>	<dbl>	<int>	<int>	<dbl>	<chr>	<int>	<chr>
1	517	515	2	830	819	11	UA	1545	N14228
2	533	529	4	850	830	20	UA	1714	N24211
3	542	540	2	923	850	33	AA	1141	N619AA
4	544	545	-1	1004	1022	-18	B6	725	N804JB
5	554	600	-6	812	837	-25	DL	461	N668DN
6	554	558	-4	740	728	12	UA	1696	N39463
7	555	600	-5	913	854	19	B6	507	N516JB
8	557	600	-3	709	723	-14	EV	5708	N829AS
9	557	600	-3	838	846	-8	B6	79	N593JB
10	550	600	2	752	715	0	AA	201	N281TAA

rename() - Change column names

```
1 flights %>% rename(tail_number = tailnum)
```

```
# A tibble: 336,776 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	1	1	517	515	2	830	819	11
2	2013	1	1	533	529	4	850	830	20
3	2013	1	1	542	540	2	923	850	33
4	2013	1	1	544	545	-1	1004	1022	-18
5	2013	1	1	554	600	-6	812	837	-25
6	2013	1	1	554	558	-4	740	728	12
7	2013	1	1	555	600	-5	913	854	19
8	2013	1	1	557	600	-3	709	723	-14
9	2013	1	1	557	600	-3	838	846	-8
10	2013	1	1	558	600	0	752	745	0

select() vs. rename()

```
1 flights %>% select(tail_number = tailnum)
```

```
# A tibble: 336,776 × 1
```

```
  tail_number  
  <chr>
```

```
1 N14228  
2 N24211  
3 N619AA  
4 N804JB  
5 N668DN  
6 N39463  
7 N516JB  
8 N829AS  
9 N593JB  
10 N3ALAA
```

```
# ... with 336,766 more rows
```

```
1 flights %>% rename(tail_number = tailnum)
```

```
# A tibble: 336,776 × 19
```

```
  year month   day dep_time sched_dep...1 dep_d...2  
  <int> <int> <int>   <int>         <int>    <dbl>  
1  2013     1     1     517           515      2  
2  2013     1     1     533           529      4  
3  2013     1     1     542           540      2  
4  2013     1     1     544           545     -1  
5  2013     1     1     554           600     -6  
6  2013     1     1     554           558     -4  
7  2013     1     1     555           600     -5  
8  2013     1     1     557           600     -3  
9  2013     1     1     557           600     -3  
10 2013     1     1     558           600     -2
```

```
# ... with 336,766 more rows, 13 more variables:
```

```
#   arr_time <int>, sched_arr_time <int>,  
#   arr_delay <dbl>, carrier <chr>, flight <int>,  
#   tail number <chr>, origin <chr>, dest <chr>.
```

pull()

```
1 names(flights)
```

```
[1] "year"          "month"          "day"            "dep_time"
[5] "sched_dep_time" "dep_delay"      "arr_time"       "sched_arr_time"
[9] "arr_delay"      "carrier"        "flight"         "tailnum"
[13] "origin"         "dest"           "air_time"       "distance"
[17] "hour"           "minute"         "time_hour"
```

```
1 flights %>% pull("year") %>% head()
```

```
[1] 2013 2013 2013 2013 2013 2013
```

```
1 flights %>% pull(1) %>% head()
```

```
[1] 2013 2013 2013 2013 2013 2013
```

```
1 flights %>% pull(-1) %>% head()
```

```
[1] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"
[3] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"
[5] "2013-01-01 06:00:00 EST" "2013-01-01 05:00:00 EST"
```

arrange() - Sort data

```
1 flights %>% filter(month==3,day==2) %>% arrange(origin, dest)
```

```
# A tibble: 765 × 19
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	3	2	1336	1329	7	1426	1432	-6
2	2013	3	2	628	629	-1	837	849	-12
3	2013	3	2	637	640	-3	903	915	-12
4	2013	3	2	743	745	-2	945	1010	-25
5	2013	3	2	857	900	-3	1117	1126	-9
6	2013	3	2	1027	1030	-3	1234	1247	-13
7	2013	3	2	1134	1145	-11	1332	1359	-27
8	2013	3	2	1412	1415	-3	1636	1630	6
9	2013	3	2	1633	1636	-3	1848	1908	-20
10	2013	3	2	1655	1700	5	1857	1924	27

arrange() & desc() - Descending order

```
1 flights %>%  
2   filter(month==3, day==2) %>%  
3   arrange(desc(origin), dest) %>%  
4   select(origin, dest, tailnum)
```

A tibble: 765 × 3

	origin	dest	tailnum
	<chr>	<chr>	<chr>
1	LGA	ATL	N928AT
2	LGA	ATL	N623DL
3	LGA	ATL	N680DA
4	LGA	ATL	N996AT
5	LGA	ATL	N510MQ
6	LGA	ATL	N663DN
7	LGA	ATL	N942DL
8	LGA	ATL	N511MQ
9	LGA	ATL	N910DE
10	LGA	ATL	N902DE

distinct() - Find unique rows

```
1 flights %>%  
2   select(origin, dest) %>%  
3   distinct() %>%  
4   arrange(origin,dest)
```

```
# A tibble: 224 × 2
```

	origin	dest
	<chr>	<chr>
1	EWR	ALB
2	EWR	ANC
3	EWR	ATL
4	EWR	AUS
5	EWR	AVL
6	EWR	BDL
7	EWR	BNA
8	EWR	BOS
9	EWR	BQN
10	EWR	DTW

mutate() - Modify / create columns

```
1 flights %>%  
2   select(year:day) %>%  
3   mutate(date = paste(year, month, day, sep="/"))
```

A tibble: 336,776 × 4

	year	month	day	date
	<int>	<int>	<int>	<chr>
1	2013	1	1	2013/1/1
2	2013	1	1	2013/1/1
3	2013	1	1	2013/1/1
4	2013	1	1	2013/1/1
5	2013	1	1	2013/1/1
6	2013	1	1	2013/1/1
7	2013	1	1	2013/1/1
8	2013	1	1	2013/1/1
9	2013	1	1	2013/1/1
10	2013	1	1	2013/1/1

summarise() - Aggregate rows

```
1 flights %>%  
2   summarize(n(), min(dep_delay), max(dep_delay))
```

A tibble: 1 × 3

	`n()`	`min(dep_delay)`	`max(dep_delay)`
	<int>	<dbl>	<dbl>
1	336776	NA	NA

```
1 flights %>%  
2   summarize(  
3     n = n(),  
4     min_dep_delay = min(dep_delay, na.rm = TRUE),  
5     max_dep_delay = max(dep_delay, na.rm = TRUE)  
6   )
```

A tibble: 1 × 3

	n	min_dep_delay	max_dep_delay
	<int>	<dbl>	<dbl>
1	336776	-43	1301

group_by()

```
1 flights %>% group_by(origin)
```

```
# A tibble: 336,776 × 19
```

```
# Groups:   origin [3]
```

	year	month	day	dep_time	sched_dep_t... ¹	dep_d... ²	arr_t... ³	sched... ⁴	arr_d... ⁵
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	2013	1	1	517	515	2	830	819	11
2	2013	1	1	533	529	4	850	830	20
3	2013	1	1	542	540	2	923	850	33
4	2013	1	1	544	545	-1	1004	1022	-18
5	2013	1	1	554	600	-6	812	837	-25
6	2013	1	1	554	558	-4	740	728	12
7	2013	1	1	555	600	-5	913	854	19
8	2013	1	1	557	600	-3	709	723	-14
9	2013	1	1	557	600	2	830	819	11

summarise() with group_by()

```
1 flights %>%
2   group_by(origin) %>%
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm = TRUE),
6     max_dep_delay = max(dep_delay, na.rm = TRUE)
7   )
```

A tibble: 3 × 4

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

Groups after summarise

```
1 flights %>%
2   group_by(origin) %>%
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE),
7     .groups = "drop_last"
8   )
```

A tibble: 3 × 4

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

```
1 flights %>%
2   group_by(origin) %>%
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE),
7     .groups = "keep"
8   )
```

A tibble: 3 × 4

Groups: origin [3]

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

count()

```
1 flights %>%
2   group_by(origin, carrier) %>%
3   summarize(
4     n = n(),
5     .groups = "drop"
6   )
```

A tibble: 35 × 3

	origin	carrier	n
	<chr>	<chr>	<int>
1	EWR	9E	1268
2	EWR	AA	3487
3	EWR	AS	714
4	EWR	B6	6557
5	EWR	DL	4342
6	EWR	EV	43939
7	EWR	MQ	2276
8	EWR	OO	6
9	EWR	UA	46087
10	EWR	US	1105

```
1 flights %>%
2   count(origin, carrier)
```

A tibble: 35 × 3

	origin	carrier	n
	<chr>	<chr>	<int>
1	EWR	9E	1268
2	EWR	AA	3487
3	EWR	AS	714
4	EWR	B6	6557
5	EWR	DL	4342
6	EWR	EV	43939
7	EWR	MQ	2276
8	EWR	OO	6
9	EWR	UA	46087
10	EWR	US	1105

mutate() with group_by()

```
1 flights %>% group_by(origin) %>%  
2   mutate(  
3     n = n(),  
4   ) %>%  
5   select(origin, n)
```

```
# A tibble: 336,776 × 2
```

```
# Groups:   origin [3]
```

	origin	n
	<chr>	<int>
1	EWR	120835
2	LGA	104662
3	JFK	111279
4	JFK	111279
5	LGA	104662
6	EWR	120835
7	EWR	120835
8	LGA	104662
9	JFK	111279

Exercises / Examples

1. How many flights to Los Angeles (LAX) did each of the legacy carriers (AA, UA, DL or US) have in May from JFK, and what was their average duration?
2. What was the shortest flight out of each airport in terms of distance? In terms of duration?
3. Which plane (check the tail number) flew out of each New York airport the most?
4. Which date should you fly on if you want to have the lowest possible average departure delay? What about arrival delay?