

Profiling & Parallelization

Lecture 20

Dr. Colin Rundel

Profiling & Benchmarking

profvis demo

```
1 n = 1e6
2 d = tibble(
3   x1 = rt(n, df = 3),
4   x2 = rt(n, df = 3),
5   x3 = rt(n, df = 3),
6   x4 = rt(n, df = 3),
7   x5 = rt(n, df = 3),
8 ) |>
9   mutate(y = -2*x1 - 1*x2 + 0*x3 + 1*x4 + 2*x5 + rnorm(n))
```

```
1 profvis::profvis(lm(y~., data=d))
```

Benchmarking - bench

```
1 d = tibble(  
2   x = runif(10000),  
3   y = runif(10000)  
4 )  
5  
6 (b = bench::mark(  
7   d[d$x > 0.5, ],  
8   d[which(d$x > 0.5), ],  
9   subset(d, x > 0.5),  
10  filter(d, x > 0.5)  
11 ))
```

A tibble: 4 × 6

expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1 d[d\$x > 0.5,]	49.4µs	56.7µs	17251.	238.11KB	46.9
2 d[which(d\$x > 0.5),]	90.7µs	103µs	9595.	268.38KB	48.3
3 subset(d, x > 0.5)	87.5µs	106.8µs	9234.	296.33KB	49.5
4 filter(d, x > 0.5)	287µs	314µs	3061.	1.48MB	18.4

Larger n

```
1 d = tibble(  
2   x = runif(1e6),  
3   y = runif(1e6)  
4 )  
5  
6 (b = bench::mark(  
7   d[d$x > 0.5, ],  
8   d[which(d$x > 0.5), ],  
9   subset(d, x > 0.5),  
10  filter(d, x > 0.5)  
11 ))
```

A tibble: 4 × 6

expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1 d[d\$x > 0.5,]	3.62ms	4.08ms	247.	13.4MB	142.
2 d[which(d\$x > 0.5),]	8.12ms	8.93ms	111.	24.8MB	140.
3 subset(d, x > 0.5)	9.46ms	10.12ms	99.0	24.8MB	120.
4 filter(d, x > 0.5)	4.81ms	5.36ms	184.	24.8MB	272.

bench - relative results

```
1 summary(b, relative=TRUE)
```

```
# A tibble: 4 × 6
```

	expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
	<bch:expr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	d[d\$x > 0.5,]	1	1	2.49	1	1.19
2	d[which(d\$x > 0.5),]	2.24	2.19	1.12	1.86	1.17
3	subset(d, x > 0.5)	2.62	2.48	1	1.86	1
4	filter(d, x > 0.5)	1.33	1.32	1.86	1.86	2.27

Parallelization

parallel

Part of the base packages in R

- tools for the forking of R processes (some functions do not work on Windows)
- Core functions:
 - `detectCores`
 - `pvec`
 - `mclapply`
 - `mcpipeline` & `mccollect`

detectCores

Surprisingly, detects the number of cores of the current system.

```
1 detectCores()
```

```
[1] 10
```

pvec

Parallelization of a vectorized function call

```
1 system.time(pvec(1:1e7, sqrt, mc.cores = 1))
```

user	system	elapsed
0.089	0.012	0.100

```
1 system.time(pvec(1:1e7, sqrt, mc.cores = 4))
```

user	system	elapsed
0.126	0.119	0.198

```
1 system.time(pvec(1:1e7, sqrt, mc.cores = 8))
```

user	system	elapsed
0.092	0.165	0.168

```
1 system.time(sqrt(1:1e7))
```

user	system	elapsed
0.018	0.018	0.038

pvec - bench::system_time

```
1 bench::system_time(pvec(1:1e7, sqrt, mc.cores = 1))
```

process	real
61.4ms	60.7ms

```
1 bench::system_time(pvec(1:1e7, sqrt, mc.cores = 4))
```

process	real
157ms	188ms

```
1 bench::system_time(pvec(1:1e7, sqrt, mc.cores = 8))
```

process	real
175ms	207ms

```
1 bench::system_time(Sys.sleep(.5))
```

process	real
57 μ s	497ms

```
1 system.time(Sys.sleep(.5))
```

user	system	elapsed
0.000	0.000	0.505

Cores by size

```
1 cores = c(1,4,8,16)
2 order = 6:8
3 f = function(x,y) {
4   system.time(
5     pvec(1:(10^y), sqrt, mc.cores = x)
6   )[3]
7 }
8
9 res = map(
10   cores,
11   function(x) {
12     map_dbl(order, f, x = x)
13   }
14 ) |>
15   do.call(rbind, args = _)
16
17 rownames(res) = paste0(cores," cores")
18 colnames(res) = paste0("10^",order)
19
20 res
```

	10^6	10^7	10^8
1 cores	0.005	0.025	0.327
4 cores	0.041	0.161	1.687
8 cores	0.047	0.174	1.361
16 cores	0.055	0.185	1.359

mclapply

Parallelized version of lapply

```
1 system.time(rnorm(1e7))
```

user	system	elapsed
0.267	0.006	0.275

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 2)))
```

user	system	elapsed
0.320	0.106	0.273

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 4)))
```

user	system	elapsed
0.326	0.100	0.179

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 8)))
```

user	system	elapsed
0.340	0.135	0.180

```
1 system.time(unlist(mclapply(1:10, function(x) rnorm(1e6), mc.cores = 10)))
```

user	system	elapsed
0.354	0.159	0.204

mcparallel

Asynchronously evaluation of an R expression in a separate process

```
1 m = mcparallel(rnorm(1e6))
2 n = mcparallel(rbeta(1e6,1,1))
3 o = mcparallel(rgamma(1e6,1,1))
4
5 str(m)
```

List of 2

```
$ pid: int 80672
$ fd : int [1:2] 4 7
- attr(*, "class")= chr [1:3] "parallelJob" "childProcess" "process"
```

```
1 str(n)
```

List of 2

```
$ pid: int 80673
$ fd : int [1:2] 5 9
- attr(*, "class")= chr [1:3] "parallelJob" "childProcess" "process"
```


mccollect

Checks `mcpaallel` objects for completion

```
1 str(mccollect(list(m,n,o)))
```

List of 3

```
$ 80672: num [1:1000000] -1.049 -1.028 0.705 -0.17 1.734 ...
```

```
$ 80673: num [1:1000000] 0.373 0.989 0.259 0.713 0.649 ...
```

```
$ 80674: num [1:1000000] 0.5921 1.2226 0.0756 2.4405 0.0465 ...
```

mccollect - waiting

```
1 p = mcparrallel(mean(rnorm(1e5)))  
2 mccollect(p, wait = FALSE, 10) # will retrieve the result (since it's fast)
```

```
$`80675`
```

```
[1] 0.001344587
```

```
1 mccollect(p, wait = FALSE)      # will signal the job as terminating
```

```
NULL
```

```
1 mccollect(p, wait = FALSE)      # there is no longer such a job
```

```
NULL
```

doMC & foreach

doMC & foreach

Packages by Revolution Analytics that provides the `foreach` function which is a parallelizable `for` loop (and then some).

- Core functions:
 - `registerDoMC`
 - `foreach, %dopar%, %do%`

registerDoMC

Primarily used to set the number of cores used by `foreach`, by default uses `options("cores")` or half the number of cores found by `detectCores` from the `parallel` package.

```
1 options("cores")
```

```
$cores
```

```
NULL
```

```
1 detectCores()
```

```
[1] 10
```

```
1 getDoParWorkers()
```

```
[1] 1
```

```
1 registerDoMC(4)
2 getDoParWorkers()
```

```
[1] 4
```

foreach

A slightly more powerful version of base `for` loops (think `for` with an `lapply` flavor). Combined with `%do%` or `%dopar%` for single or multicore execution.

```
1 for(i in 1:10) {  
2   sqrt(i)  
3 }  
4  
5 foreach(i = 1:5) %do% {  
6   sqrt(i)  
7 }
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] 1.414214
```

```
[[3]]
```

```
[1] 1.732051
```

```
[[ 4 ]]
```

```
[ 1]  2
```

foreach - iterators

`foreach` can iterate across more than one value, but it doesn't do length coercion

```
1 foreach(i = 1:5, j = 1:5) %do% {  
2   sqrt(i^2+j^2)  
3 }
```

```
[[1]]  
[1] 1.414214
```

```
[[2]]  
[1] 2.828427
```

```
[[3]]  
[1] 4.242641
```

```
[[4]]  
[1] 5.656854
```

```
[[5]]
```

```
1 foreach(i = 1:5, j = 1:2) %do% {  
2   sqrt(i^2+j^2)  
3 }
```

```
[[1]]  
[1] 1.414214
```

```
[[2]]  
[1] 2.828427
```


foreach - combining results

```
1 foreach(i = 1:5, .combine='c') %do% {  
2   sqrt(i)  
3 }
```

```
[1] 1.000000 1.414214 1.732051 2.000000 2.236068
```

```
1 foreach(i = 1:5, .combine='cbind') %do% {  
2   sqrt(i)  
3 }
```

```
      result.1 result.2 result.3 result.4 result.5  
[1,]          1 1.414214 1.732051          2 2.236068
```

```
1 foreach(i = 1:5, .combine='+') %do% {  
2   sqrt(i)  
3 }
```

```
[1] 8.382332
```

foreach - parallelization

Swapping out `%do%` for `%dopar%` will use the parallel backend.

```
1 registerDoMC(4)
```

```
1 system.time(foreach(i = 1:10) %dopar% mean(rnorm(1e6))))
```

user	system	elapsed
0.296	0.027	0.110

```
1 registerDoMC(8)
```

```
2 system.time(foreach(i = 1:10) %dopar% mean(rnorm(1e6))))
```

user	system	elapsed
0.317	0.045	0.082

```
1 registerDoMC(12)
```

```
2 system.time(foreach(i = 1:10) %dopar% mean(rnorm(1e6))))
```

user	system	elapsed
0.329	0.047	0.074



furrr / future

```
1 system.time( purrr::map(c(1,1,1), Sys.sleep) )
```

user	system	elapsed
0.001	0.000	3.012

```
1 system.time( furrr::future_map(c(1,1,1), Sys.sleep) )
```

user	system	elapsed
0.050	0.007	3.083

```
1 future::plan(future::multisession) # See also future::multicore  
2 system.time( furrr::future_map(c(1,1,1), Sys.sleep) )
```


user	system	elapsed
0.163	0.004	1.454

Example - Bootstrapping

Bootstrapping is a resampling scheme where the original data is repeatedly reconstructed by taking a samples of size n (with replacement) from the original data, and using that to repeat an analysis procedure of interest. Below is an example of fitting a local regression (`loess`) to some synthetic data, we will construct a bootstrap prediction interval for this model.

```
1  set.seed(3212016)
2  d = data.frame(x = 1:120) |>
3    mutate(y = sin(2*pi*x/120) + runif(length(x), -1, 1))
4
5  l = loess(y ~ x, data=d)
6  p = predict(l, se=TRUE)
7
8  d = d |> mutate(
9    pred_y = p$fit,
10   pred_y_se = p$se.fit
11 )
```

```
1 ggplot(d, aes(x,y)) +  
2   geom_point(color="gray50") +  
3   geom_ribbon(  
4     aes(ymin = pred_y - 1.96 * pred_y_se,  
5         ymax = pred_y + 1.96 * pred_y_se),  
6     fill="red", alpha=0.25  
7   ) +  
8   geom_line(aes(y=pred_y)) +  
9   theme_bw()
```



What to use when?

Optimal use of multiple cores is hard, there isn't one best solution

- Don't underestimate the overhead cost
- Experimentation is key
- Measure it or it didn't happen
- Be aware of the trade off between developer time and run time

BLAS and LAPACK

Statistics and Linear Algebra

An awful lot of statistics is at its core linear algebra.

For example:

- Linear regression models, find

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

- Principle component analysis
 - Find $T = XW$ where W is a matrix whose columns are the eigenvectors of $X^T X$.
 - Often solved via SVD - Let $X = U\Sigma W^T$ then $T = U\Sigma$.

Numerical Linear Algebra

Not unique to Statistics, these are the type of problems that come up across all areas of numerical computing.

- Numerical linear algebra \neq mathematical linear algebra
- Efficiency and stability of numerical algorithms matter
 - Designing and implementing these algorithms is hard
- Don't reinvent the wheel - common core linear algebra tools (well defined API)

BLAS and LAPACK

Low level algorithms for common linear algebra operations

BLAS

- **B**asic **L**inear **A**lgebra **S**ubprograms
- Copying, scaling, multiplying vectors and matrices
- Origins go back to 1979, written in Fortran

LAPACK

- **L**inear **A**lgebra **P**ackage
- Higher level functionality building on BLAS.
- Linear solvers, eigenvalues, and matrix decompositions
- Origins go back to 1992, mostly Fortran (expanded on LINPACK, EISPACK)

Modern variants?

Most default BLAS and LAPACK implementations (like R's defaults) are somewhat dated

- Written in Fortran and designed for a single cpu core
- Certain (potentially non-optimal) hard coded defaults (e.g. block size).

Multithreaded alternatives:


- ATLAS - Automatically Tuned Linear Algebra Software
- OpenBLAS - fork of GotoBLAS from TACC at UTexas
- Intel MKL - Math Kernel Library, part of Intel's commercial compiler tools
- cuBLAS / Magma - GPU libraries from Nvidia and UTK respectively
- Accelerate / vecLib - Apple's framework for GPU and multicore computing

OpenBLAS Matrix Multiply Performance

```
1 x=matrix(runif(5000^2),ncol=5000)
2
3 sizes = c(100,500,1000,2000,3000,4000,5000)
4 cores = c(1,2,4,8,16)
5
6 sapply(
7   cores,
8   function(n_cores)
9   {
10     flexiblas::flexiblas_set_num_threads(n_cores)
11     sapply(
12       sizes,
13       function(s)
14       {
15         y = x[1:s,1:s]
16         system.time(y %*% y)[3]
17       }
18     )
19   }
20 )
```

n	1 core	2 cores	4 cores	8 cores	16 cores
100	0.000	0.000	0.000	0.000	0.000
500	0.004	0.003	0.002	0.002	0.004
1000	0.028	0.016	0.010	0.007	0.009
2000	0.207	0.110	0.058	0.035	0.039
3000	0.679	0.352	0.183	0.103	0.081
4000	1.587	0.816	0.418	0.227	0.145
5000	3.104	1.583	0.807	0.453	0.266

Matrix Multiply of (n x n) matrices



Matrix Multiply of (n x n) matrices

